Virtual Machine

Memory Management I

資工三 B00902064 宋昊恩
資工三 B00902110 余孟桓

Outline

- Brief introduction
- Motivation
- Terminology
- Memory Demand Detection
- Memory Reclamation
- Memory Sharing
- Memory Utilization and Isolation
- Conclusion

Brief Introduction

Brief introduction

- Memory Resource Management in VMWare ESX server [1]
  - C. A. Waldspurger, OSDI '02: Proceedings of the 5th symposium on Operating systems design and implementation, 2002
- Collaborative Memory Management in Hosted Linux Environments [2]
  - Martin Schwidefsky, et. al., Linux Symposium, 2006
- Satori: Enlighted Page Sharing [3]
  - G. Milos, et. al., USENIX 2009
- Dynamic memory balancing for virtual machines [4]
  - Weiming Zhao, et. al., ACM SIGOPS Operating Systems Review, Volume 43 Issue 3, July 2009

Motivation

Motivation

- Two solution direction in terms of mechanism:
  - Memory Reclamation
  - Memory Sharing
- Two solution direction in terms of emphasis phase:
  - Memory Utilization
  - Memory Isolation
- Two solution direction in terms of fulfillment:
  - Modification on guest OS or not

Terminology

Terminology

- Overcommitment
  - Total size configured for all running VM exceeds the total amount of actual machine memory.

- Shadow Page Table
  - A mechanism for preserving the relationship between virtual address and machine address.
  - TLB then cache the mapping from Shadow Page Table.

Terminology (cont.)

- COW
  - Copy-on-write, a mechanism to save memory space. When a minor-refined copy is created, it is unnecessary to duplicate whole data, but only the differences between them.

- Paravirtualization (enlightenment, used in [3])
  - Different from full-virtualization, each guest OS system is aware of one another. In this case, the entire system can work together as a cohesive unit.

Memory Demand Detection

Memory Demand Detection

- Motivation:
  - To prevent VM idle and resources waste, it is important to get clear memory demand for each VM.

- Solution:
  - Sampling VM for specific span of time. [1]
  - LRU-based statistics. [4]

Sampling VM for specific span of time

- Description:
  - Determine span of time T, then invalidate N physical pages related TLB and MMU state in uniform distribution.
  - Once guest trying to re-establish mappings, counter increases.

- Drawbacks:
  - Tradeoff between estimation accuracy and overhead.
  - Cannot tell the VM performance when more or less memory is allocated.

LRU-based statistics

- Description:
  - Every page are cold initially, all accessibilities are removed.
  - When traps, it turns hot, accessibility is recovered.
  - When there are too many hot pages, LRU one will turn cold.
  - Trap only cold one, will not cause too much overhead.

- Drawbacks:
  - Cannot detect swap usage.
  - Too many pages to preserve.
  - Updating needs linear search.

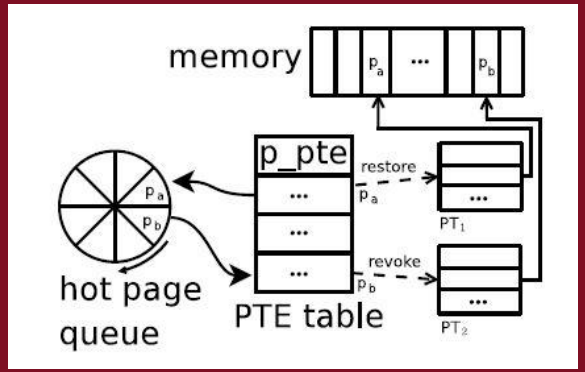

LRU-based statistics (cont.)

- Cannot detect swap usage
  - Add another background process to collect information, and send back to VMM.
- Too many pages to preserve
  - Collect G consecutive pages as a node.
- Updating needs linear search
  - Assume good locality exists.

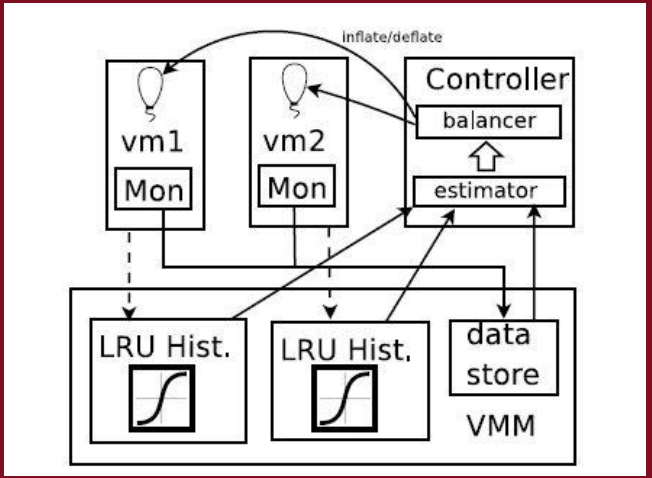

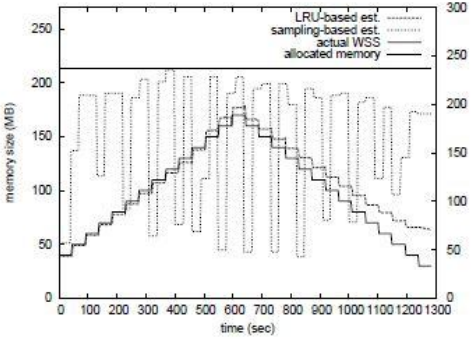

(a) `mono` (40 MB to 170 MB)

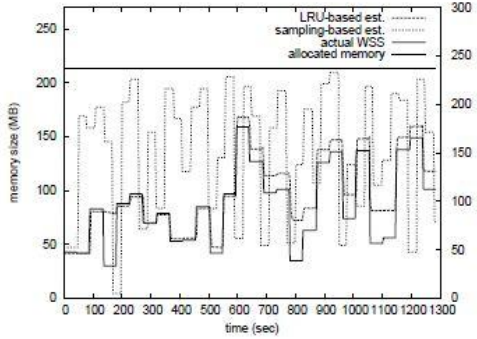

(b) `random` (40 MB to 170 MB)

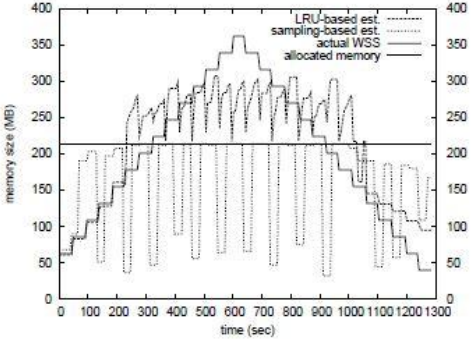

(c) `mono` (40 MB to 350 MB)

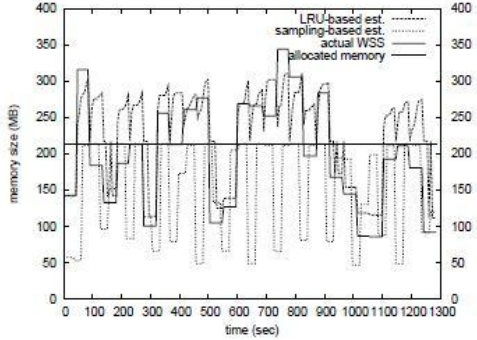

(d) `random` (40 MB to 350 MB)

Memory Reclamation

Memory Reclamation

- Motivation:
  - Commodity operating system don’t support dynamic changes to physical memory sizes.
  - But we still want to improve memory utilities.

- Solutions:
  - Page Replacement
  - Improved Page Replacement [1]
  - Ballooning [1]
  - MEmory Balancer (MEB) [4]

Page Replacement

- Description:
  - Host level pages

- Drawbacks:
  - Host-level knows little about guest OS memory situation.
  - Double Paging Problem
    - Page swapped out by host-level before swapping out by guest OS.
    - It can be improved by Randomized page replacement strategy (used by ESX server).

Ballooning

- Description:
  - A module loaded into guest OS as a pseudo-device driver.
  - When facing memory pressure, balloon inflates and notify VMM to reclaim physical pages it gain.

- Drawbacks:
  - It may be uninstalled, disabled explicitly.
  - It is not available while OS boosting.
  - Cannot reclaim memory quickly.
  - Each balloon has a minimum allocation.

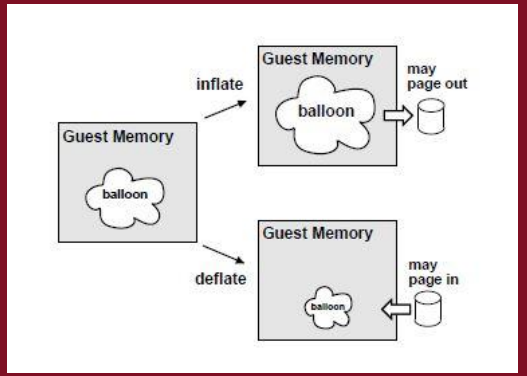

Collaborative memory management (CMM)

- Collaborative memory management (CMM)
  - Infrequent Ballooning (CMM1)
    - Apply sufficient long term request
  - Page Replacement (CMM2)
    - Guest OS maintain *page status*
    - Host OS maintain *page resident*
  - Used by zSeries and z/VM

MEmory Balancer (MEB)

- Description:
  - Based on LRU statistics, it do dynamic memory resizing.
  - If the sum of VM memory demand can be fulfilled, then done.
  - Else, a delicate optimized algorithm is needed.

Memory Sharing

Memory Sharing

- Motivation:
  - Besides adjusting memory allocation between VMs, many VM shares similar process on same OS platform.
  - Prevent malicious virtual machine take advantage of memory reclamation.

- Solutions:
  - Transparent Page Sharing
  - Content-based Page sharing [1]
  - Enlightened page sharing [3]

Transparent Page Sharing

- Description:
  - Introduced by Disco, one copies are identified, multiple guest physical pages are mapped to the same machine page.
- Drawbacks:
  - It needs several guest OS modifications to identify copies.
    - For example, it hooked some code to routine “bcopy”.

Content-based Page Sharing

- Description:
  - Initially, all pages are ordinary pages, and are hashed.
  - Once hash value meets, recalculate hash value.
    - If changed, remove old page and hash value.
    - If not changed, mark the page as COW.
  - It need to scan pages frequently.

- Drawbacks:
  - Tradeoff between scanning frequency and overhead.
  - Hard to discover short-lived share memory. (< 40 min)

Content-based Page Sharing (cont.)

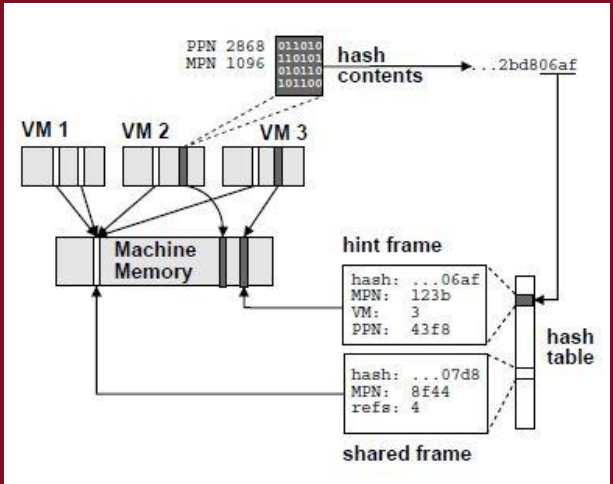

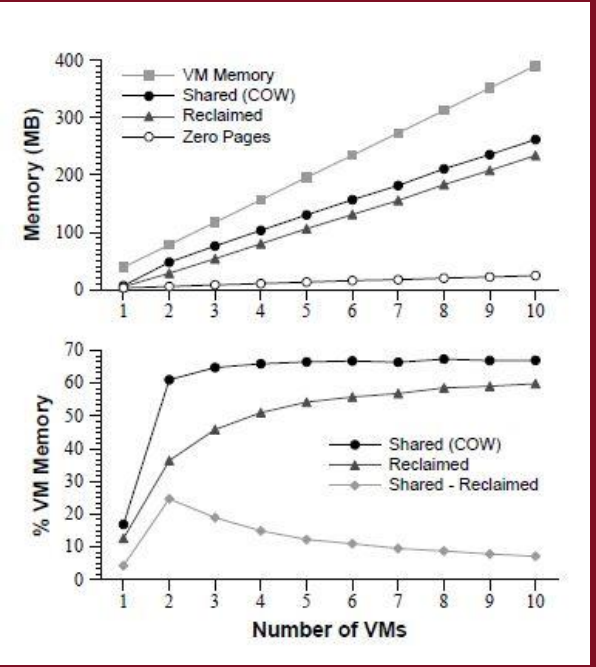

Enlighted Page Sharing

- Description:
  - Modify virtual disk subsystem, and implement Sharing-aware Block Devices.
  - Detect sharing directly when data is read from disk.

- Advantages:
  - Avoid scanning overhead.
  - Detect short-lived sharing immediately.

Enlighted Page Sharing (Cont.)

- Drawbacks:
  - Cannot detect consequent memory writes.
  - Lots of modification, including hypervisor modification, sharing-aware block device addition, and adding repayment FIFO to guest OS kernel.

Memory Utilization and Isolation

Memory Utilization

- Motivation:
  - No matter how many resources one claim, the top principle is maximize the efficiency of memory usage.
- Advantages:
  - Avoid resources idle.
- Drawbacks:
  - Some malicious VM may utilize auto memory balancing mechanism to gain unreasonable amount of resources.

Memory Isolation

- Motivation:
  - No matter how imbalancing memory usage are, VM should follows isolation principles, and not affected by others.
- Advantages:
  - Can prevent malicious user, and preserve one's privilege.
- Drawbacks:
  - Most of the time, malicious user does not exist.
  - Large amount of memory are wasted.

Some mechanisms

- Min-funding revocation:
  - One paid more money, the least valuable one becomes victim.

- Page-share advantage:
  - One share more pages with others uses more memory space.
  - No matter one page is shared or changed, only VMs sharing that certain page are involved.

Some mechanisms (cont.)

- Idle memory tax:
  - $\rho = S / (P \cdot (f + k \cdot (1 - f)))$
    - ρ : shares-per-page ratio
    - S : shares
    - P : number of pages
    - $f = t / n$: active rate
    - $k = 1 / (1 - \tau)$: idle page cost
    - τ : tax rate
    - t : touched page account
    - n : number of random sample pages

Conclusion

Conclusion

- There are always diverse choices for us to choose.
  - It is hard to balance between accuracy and overhead.
  - It is hard to balance between performance and isolation.
  - It is hard to determine whether revised OS is needed or not.