

Big Data Final Project Report

Titanic: Survival Prediction

This is a general statment for this project. It means nothing actually. What I want to do is just to increase nonsenses so that I can test whether genBlog.py works normally or not.

Teammate

- R02922164 邵 飛
- B00902064 宋昊恩
- B00902048 吳瑞洋
- B00902042 詹舜傑

Data Information

Input

- Features
 - Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, Embarked.
- Features Meaning
 - Pclass: class of accommodation.
 - SibSp: number of sibling and spouse on board.
 - Parch: number of parents and children on board.
 - Ticket: ticket id.
 - Fare: fare paid.
 - Cabin: cabin accomodated.
 - Embarked: port of embarkation.

Output

- Single Label
 - Survived or Not Survived (1/0).

Data Size

- Number of Instance
 - Train: 891
 - Test: 418
- Number of Feature
 - Numerical: 4
 - Categorical: 3
 - Nominated: 3

Onboard Evaluation

- Accuracy (TP/TP+FP)

Tools and Model Selection

Tools

- Pandas: Python package, used for *data manipulation*
- Sklearn: Python package, used for *data mining* and *data analysis*

Linear Model


Logistic Regression

- Model Introduction
 - One of the linear models, which is widely used to solve machine learning problems

- Formula:

$$\min_{\mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^l \log(1 + e^{-y_i \mathbf{w}^T \mathbf{x}_i}).$$

- Figure:




Linear Support Vector Machine

- Model Introduction
 - Model will try to find out a hyperplane to separate data points in the space spanned by features.
 - Formula:

$$\min_{\mathbf{w}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^l \xi(\mathbf{w}; \mathbf{x}_i, y_i)$$

◦ Figure:




Kernel Model

Support Vector Machine

- Model Introduction
 - Model will use RBF kernel to map data points into space with infinite dimension, then try to find out a hyperplane to separate data points.
 - Its performance should cover *Linear SVC*.
 - Formula:

$$\begin{aligned} \min_{w, b, \xi} \quad & \frac{1}{2} w^T w + C \sum_{i=1}^l \xi_i \\ \text{subject to} \quad & y_i (w^T \phi(x_i) + b) \geq 1 - \xi_i, \\ & \xi_i \geq 0, i = 1, \dots, l, \end{aligned}$$

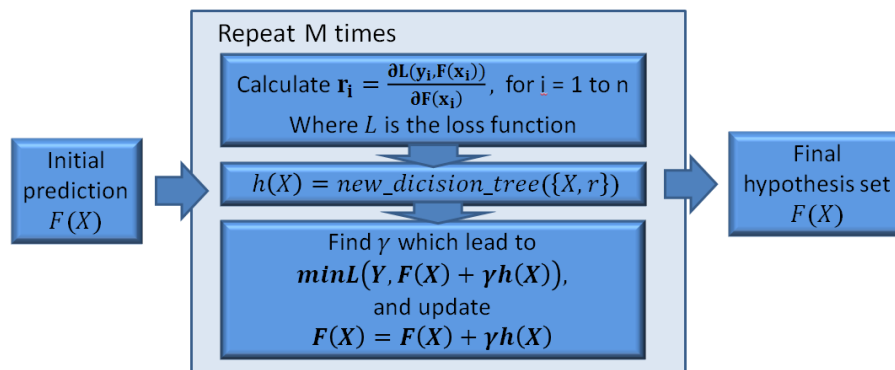
- Figure: (mapping into a space with higher dimension)




Tree-based Model

Gradient Boosting Classifier

- Model Introduction
 - Tree-based model with gradient descent update
 - Formula:



- Figure:




Random Forest Classifier

- Model Introduction
 - Tree-based model, ensemble with many out-of-bag decision trees

- Formula:

$$\hat{y} = \frac{1}{m} \sum_{j=1}^m \sum_{i=1}^n W_j(x_i, x') y_i = \sum_{i=1}^n \left(\frac{1}{m} \sum_{j=1}^m W_j(x_i, x') \right) y_i$$

- Figure:



AdaBoost Classifier

- Model Instruction
 - Selects only those features known to improve the predictive power of the model
 - Formula:
 - (a) Train classifier with respect to the weighted sample set $\{S, \mathbf{d}^{(t)}\}$ and obtain hypothesis $h_t : \mathbf{x} \mapsto \{-1, +1\}$, i.e. $h_t = \mathcal{L}(S, \mathbf{d}^{(t)})$
 - (b) Calculate the weighted training error ε_t of h_t :

$$\varepsilon_t = \sum_{n=1}^N d_n^{(t)} \mathbf{I}(y_n \neq h_t(\mathbf{x}_n)) ,$$


- (c) Set:

$$\alpha_t = \frac{1}{2} \log \frac{1 - \varepsilon_t}{\varepsilon_t}$$

- (d) Update weights:

$$d_n^{(t+1)} = d_n^{(t)} \exp \{-\alpha_t y_n h_t(\mathbf{x}_n)\} / Z_t ,$$

◦ Figure:



Feature Engineering

- Numerical
 - Features
 - Age, SibSp, Parch, Fare
 - Preprocessing
 - Impute NA with mean value
 - Standard-scaling
- Categorical
 - Features
 - Pclass, Sex, Embarked
 - Preprocessing
 - No NA is discovered
 - Binary-feature Expansion
- Nominated
 - Features
 - Name, Ticket, Cabin
 - Preprocessing
 - Lots of NA value (ex: more than 90% NA in *Cabin*)
 - Hard to use without adding human knowledge (ex: *Name*)
 - We just eliminate them in this step

Off-board Experiment Design

- Since there are few data for this problem, we must have a robust way to prevent overfitting. Then, we just apply 5-fold cross-validation for all model evaluation.
- Though we are really careful about the overfitting problem, we still find out that there are 0.04 percent difference in accuracy between off-board and on-board.

Model Performance Comparison

Linear Model

Logistic Regression

- Best Parameters $C=10$, $\text{random_state}=514$
- Performance

	Train	Test
Valid	0.80387	0.70020
Board		0.76555

Linear SVC

- Best parameters $C=10$, $\text{random_state}=514$
- Performance

	Train	Test
Valid	0.70078	0.79460
Board		0.75598

Kernel Model

Support Vector Machine

- Best Parameters: C=1, gamma=0.125, random_state=514
- Performance

	Train	Test
Valid	0.80387	0.70021
Board		0.76555

Tree-based Model

Gradient Boosting Classifier

- Best Parameters estimator=500, depth=5, random_state=514
- Performance

	Train	Test
Valid	0.89870	0.82041
Board		0.77990

Random Forest Classifier

- Best Parameters: estimator=20, depth=5, random_state=514
- Performance

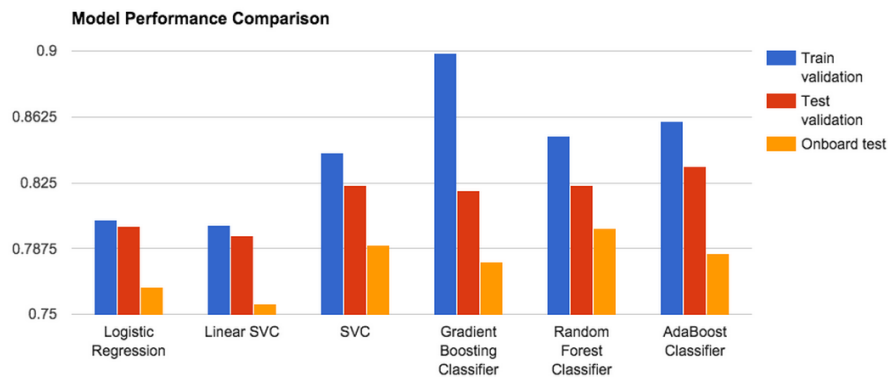
	Train	Test
Valid	0.85156	0.82378
Board		0.79904

AdaBoost Classifier

- Best Parameters estimator=30, depth=3, learning_rate= 0.2
- Performance

	Train	Test
Valid	0.85972	0.83438
Board		0.78469

Comparison from Figure



Model Ensemble

- We choose the best answer collected from each model, including SVC, GBM, Random Forest and Adaboost, and aggregate them to gain on-board score *0.79904*, which is exactly the same as the *Random Forest* one.
- One possible reason is that there is nearly nothing further can be learn from our current feature set, so different models have almost the same answer.
- To have advanced score, we can either put more efforts on nominated features or try robust feature selection for each model to enhance the model exclusiveness.

Conclusion

- We implement six ML models in this *Titanic* problem and get 0.79904 as our best result. There are several points we learn from this competition, listed as follows:
 - Some ML models have similar performance on one ML problem, i.e. Tree-based models.
 - Though some people make use of the well-known knowledge to gain 100 percent performance, this is not our main purpose in this competition. We just try to make use of what we have learned in this course.
 - There is a consistent gap between off-board and on-board score for all models. This may be caused by the imbalanced sampling in official data.

Reference

- Python Package: Scikit Learn <http://scikit-learn.org/stable/>
- Python Package: Pandas <http://pandas.pydata.org/>
- Python Software for Convex Optimization - Documentation <http://cvxopt.org/>
- A Library for Large Linear Classification <http://www.csie.ntu.edu.tw/~cjlin/papers/liblinear.pdf>
- A Library for Support Vector Machine <http://www.csie.ntu.edu.tw/~cjlin/papers/libsvm.pdf>
- Wiki page for Support Vector Machine https://en.wikipedia.org/wiki/Support_vector_machine
- Wiki page for Gradient Boosting https://en.wikipedia.org/wiki/Gradient_boosting
- Wiki page for Random Forest https://en.wikipedia.org/wiki/Random_forest
- Wiki page for AdaBoost <https://en.wikipedia.org/wiki/AdaBoost>